22/Nov/2023

Consistent Hashing

Stateful LB

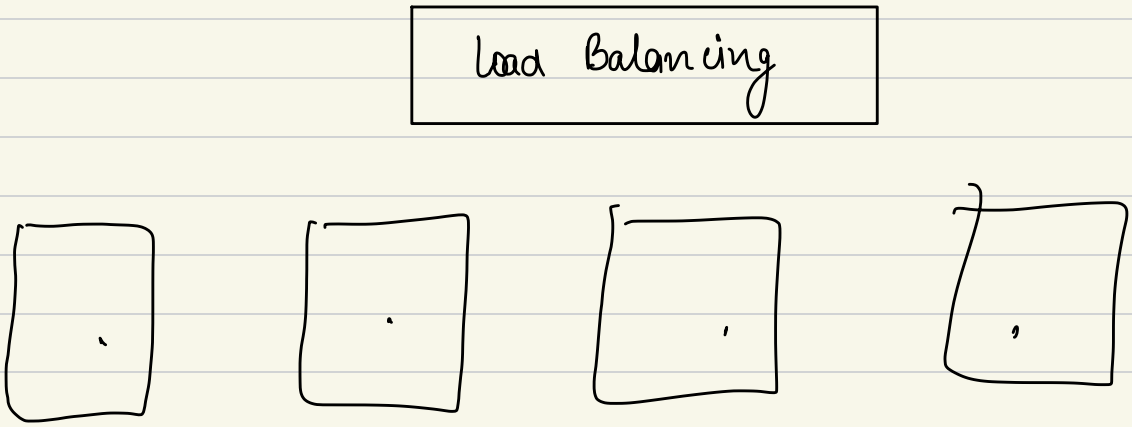

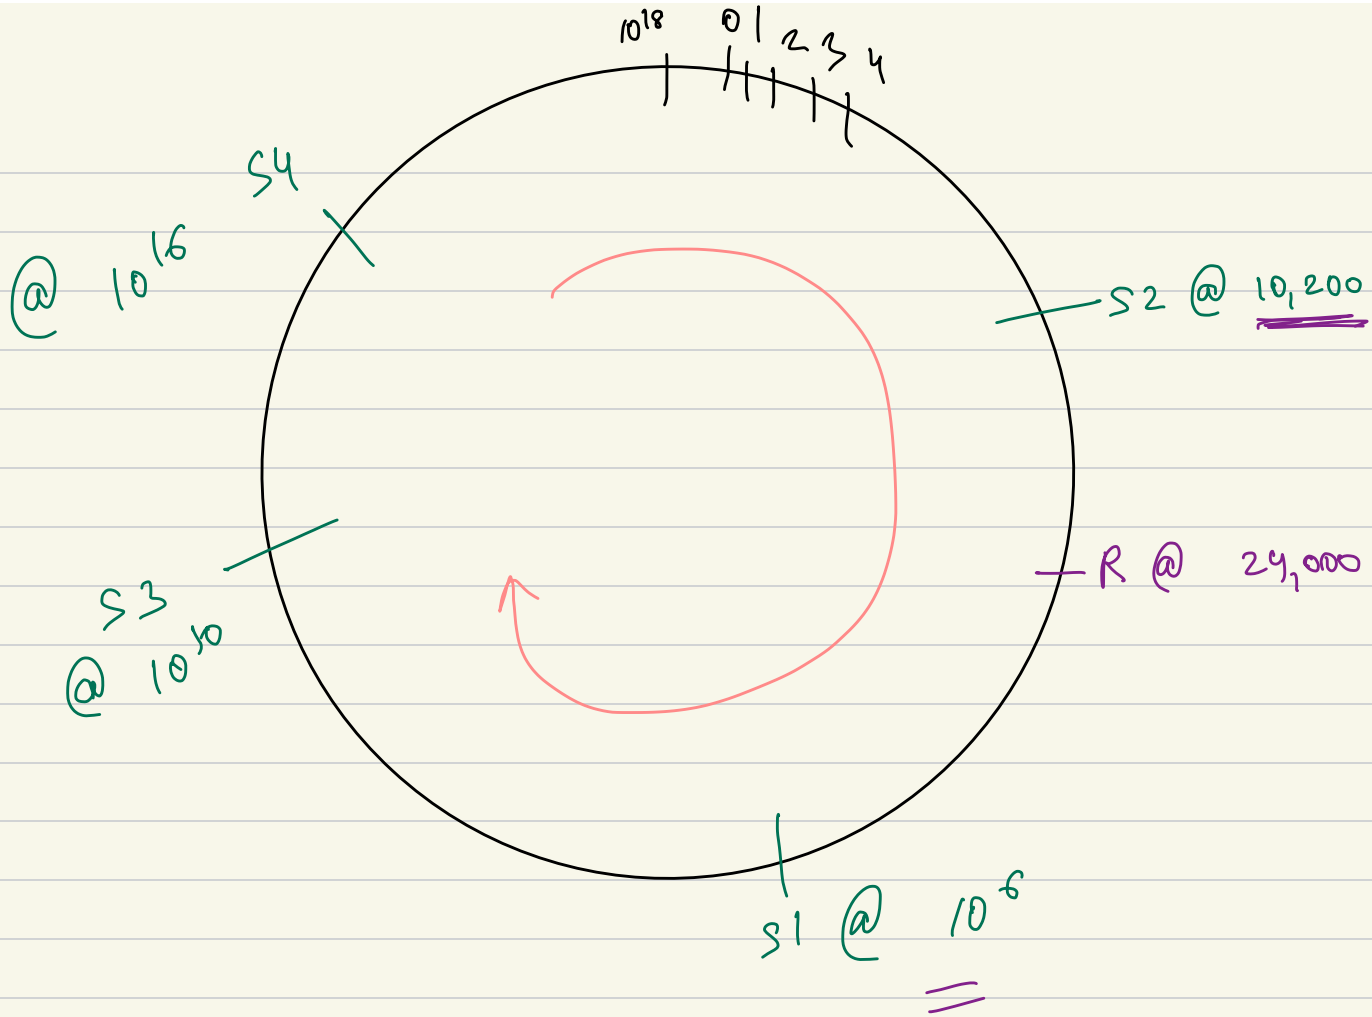

sorted array of server positions

A =

10,200	10^6	10^{10}	10^{16}
S2	S1	S3	S4

Rid=107

Load Balancer

Step 1 : $H_R(107) = 24,000$

Step 2 : Binary Search (A) to find ceiling

log(|S|) time

PROBLEM:

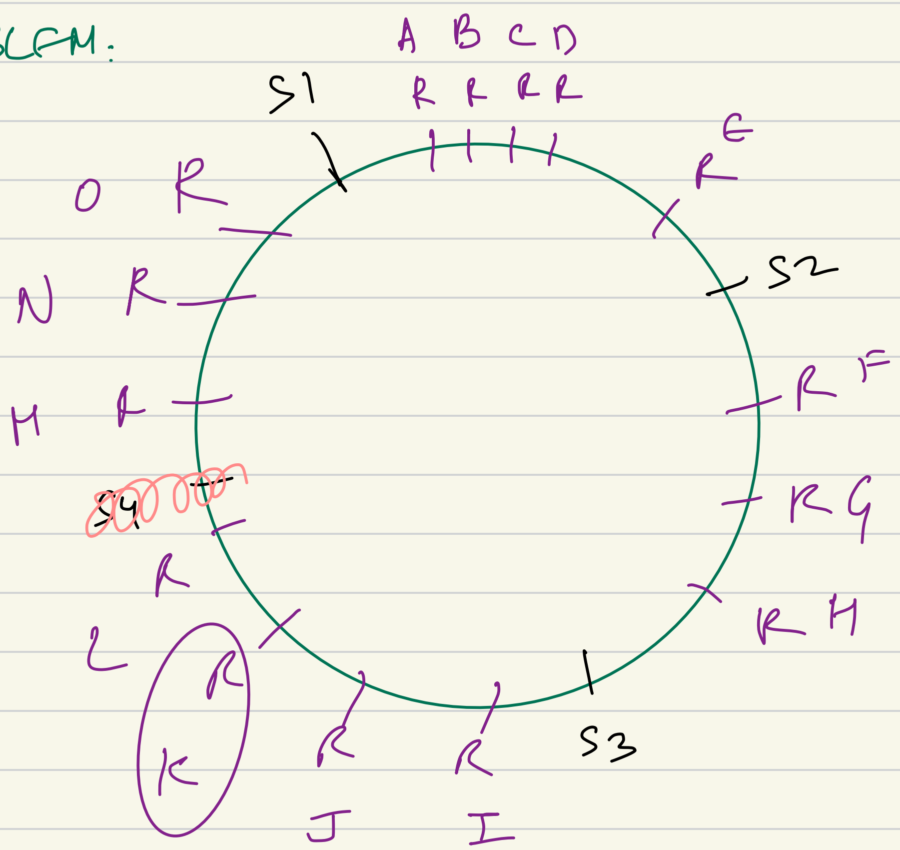

S2 ↔

S 3 ↔

S 1 ↑↑

Now S4 dies

With this approach of C.N, we have solved the problem of asking every state to move,

BUT.TTT.---

We are being unfair with re-distribution

Now S1 dies 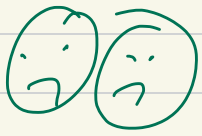

S2 ↑↑↑

S3 ⟷

Cascading failure 

(domino falling ripple failure)

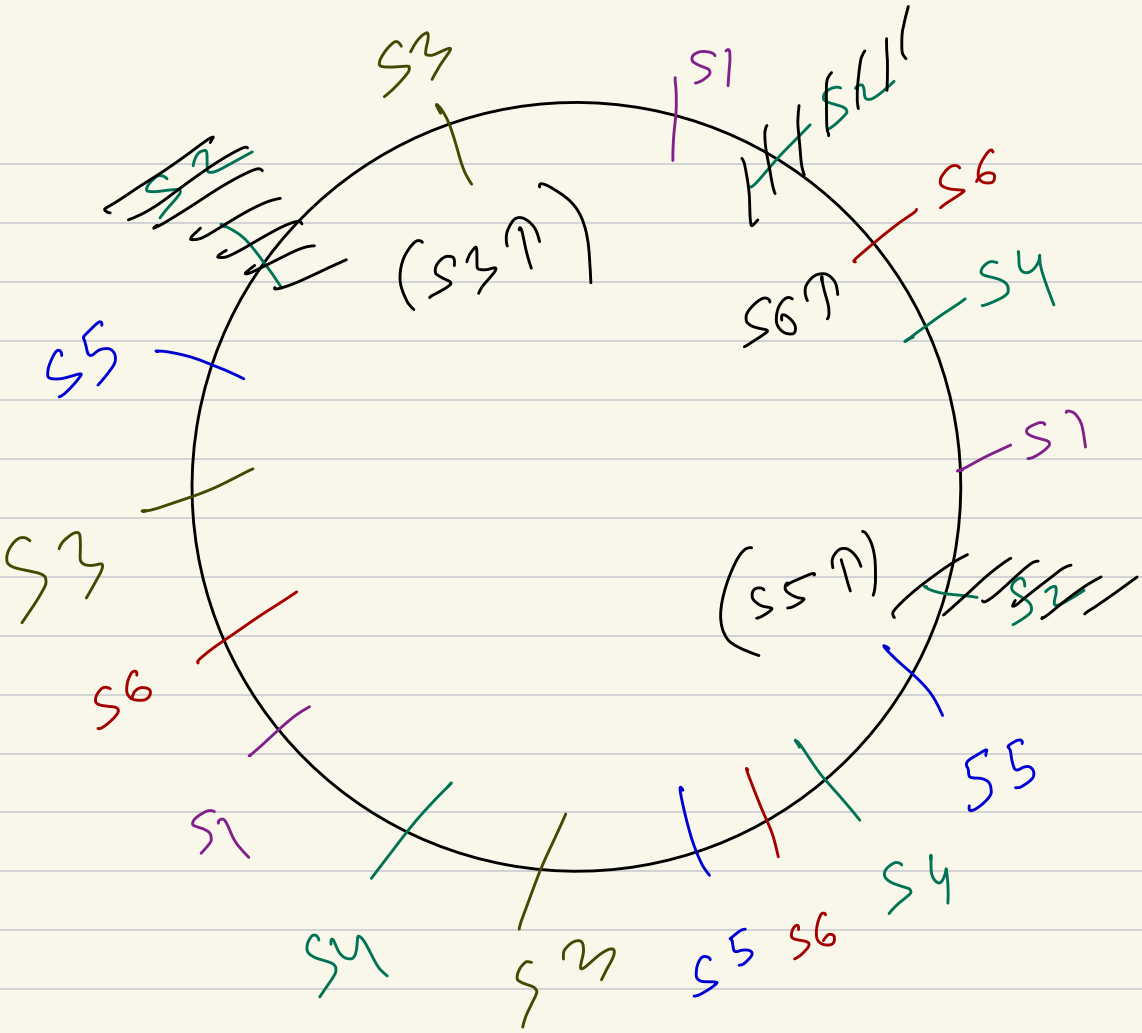

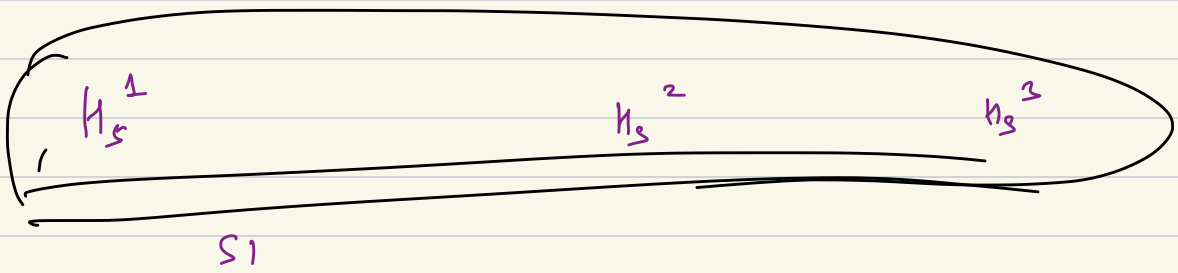

$H_s^1(s_1) = \text{value}$

$H_s^2(s_1) = \text{value } 2$

$H_s^3(s_1) = \text{value } 3$

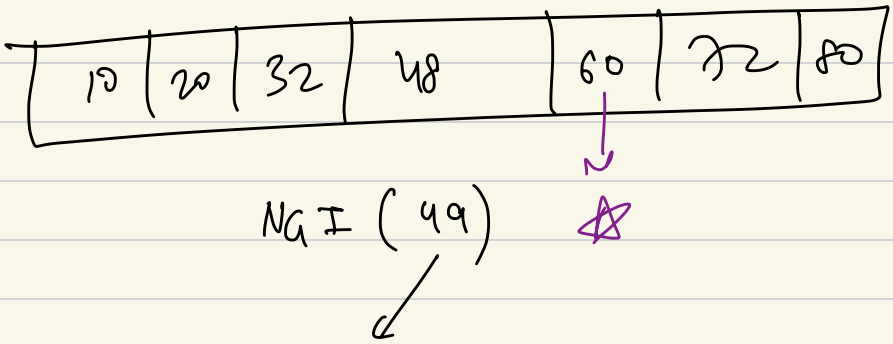

$H_R(R_{107}) = 49$

C.H Redis

Rafka

etc

etc

$H_R(R_{111}) = 950$

NGI ⇒ S1

$H_R(R_{107}) = 650$

A =

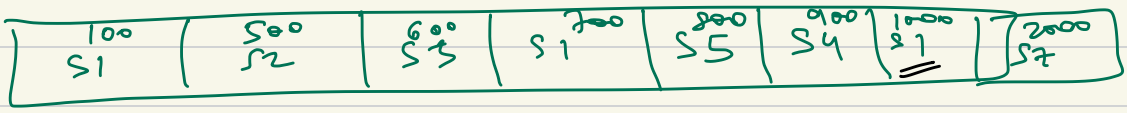

S1 dies

$H_R(R_{107}) = 650$ NGI ⇒ S5

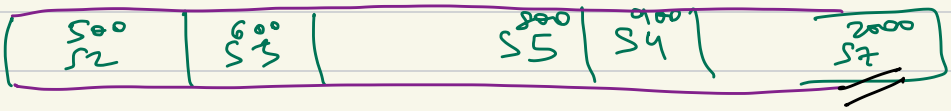

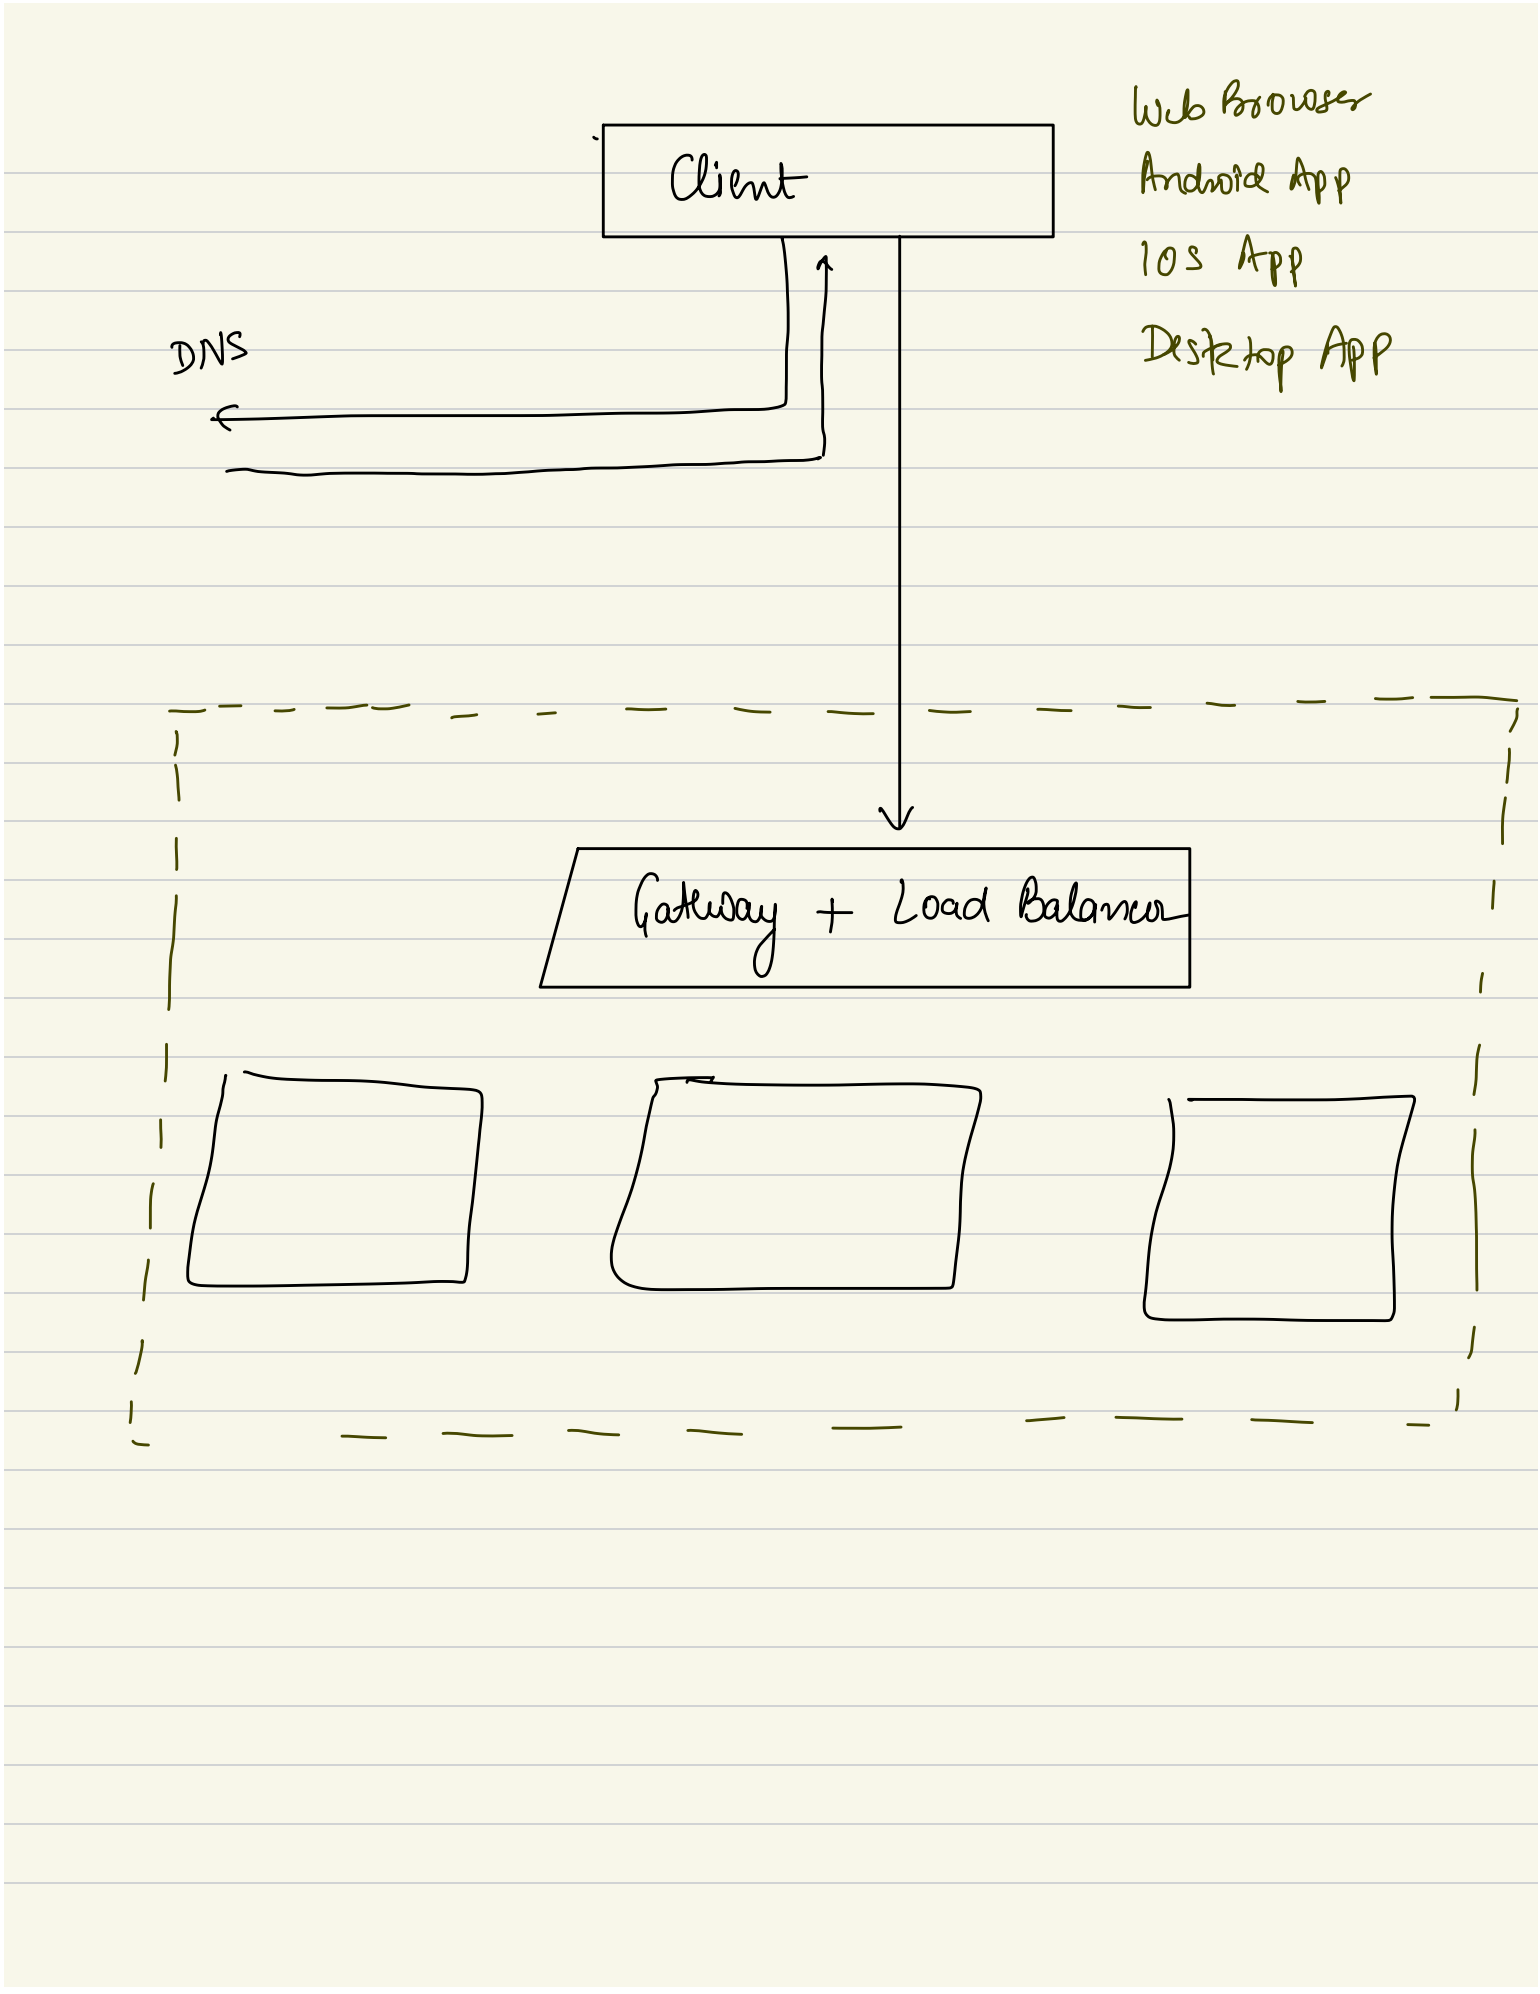
Client
Web Browser
Android App
IOS App
Desktop App
DNS
Gateway + Load Balancer

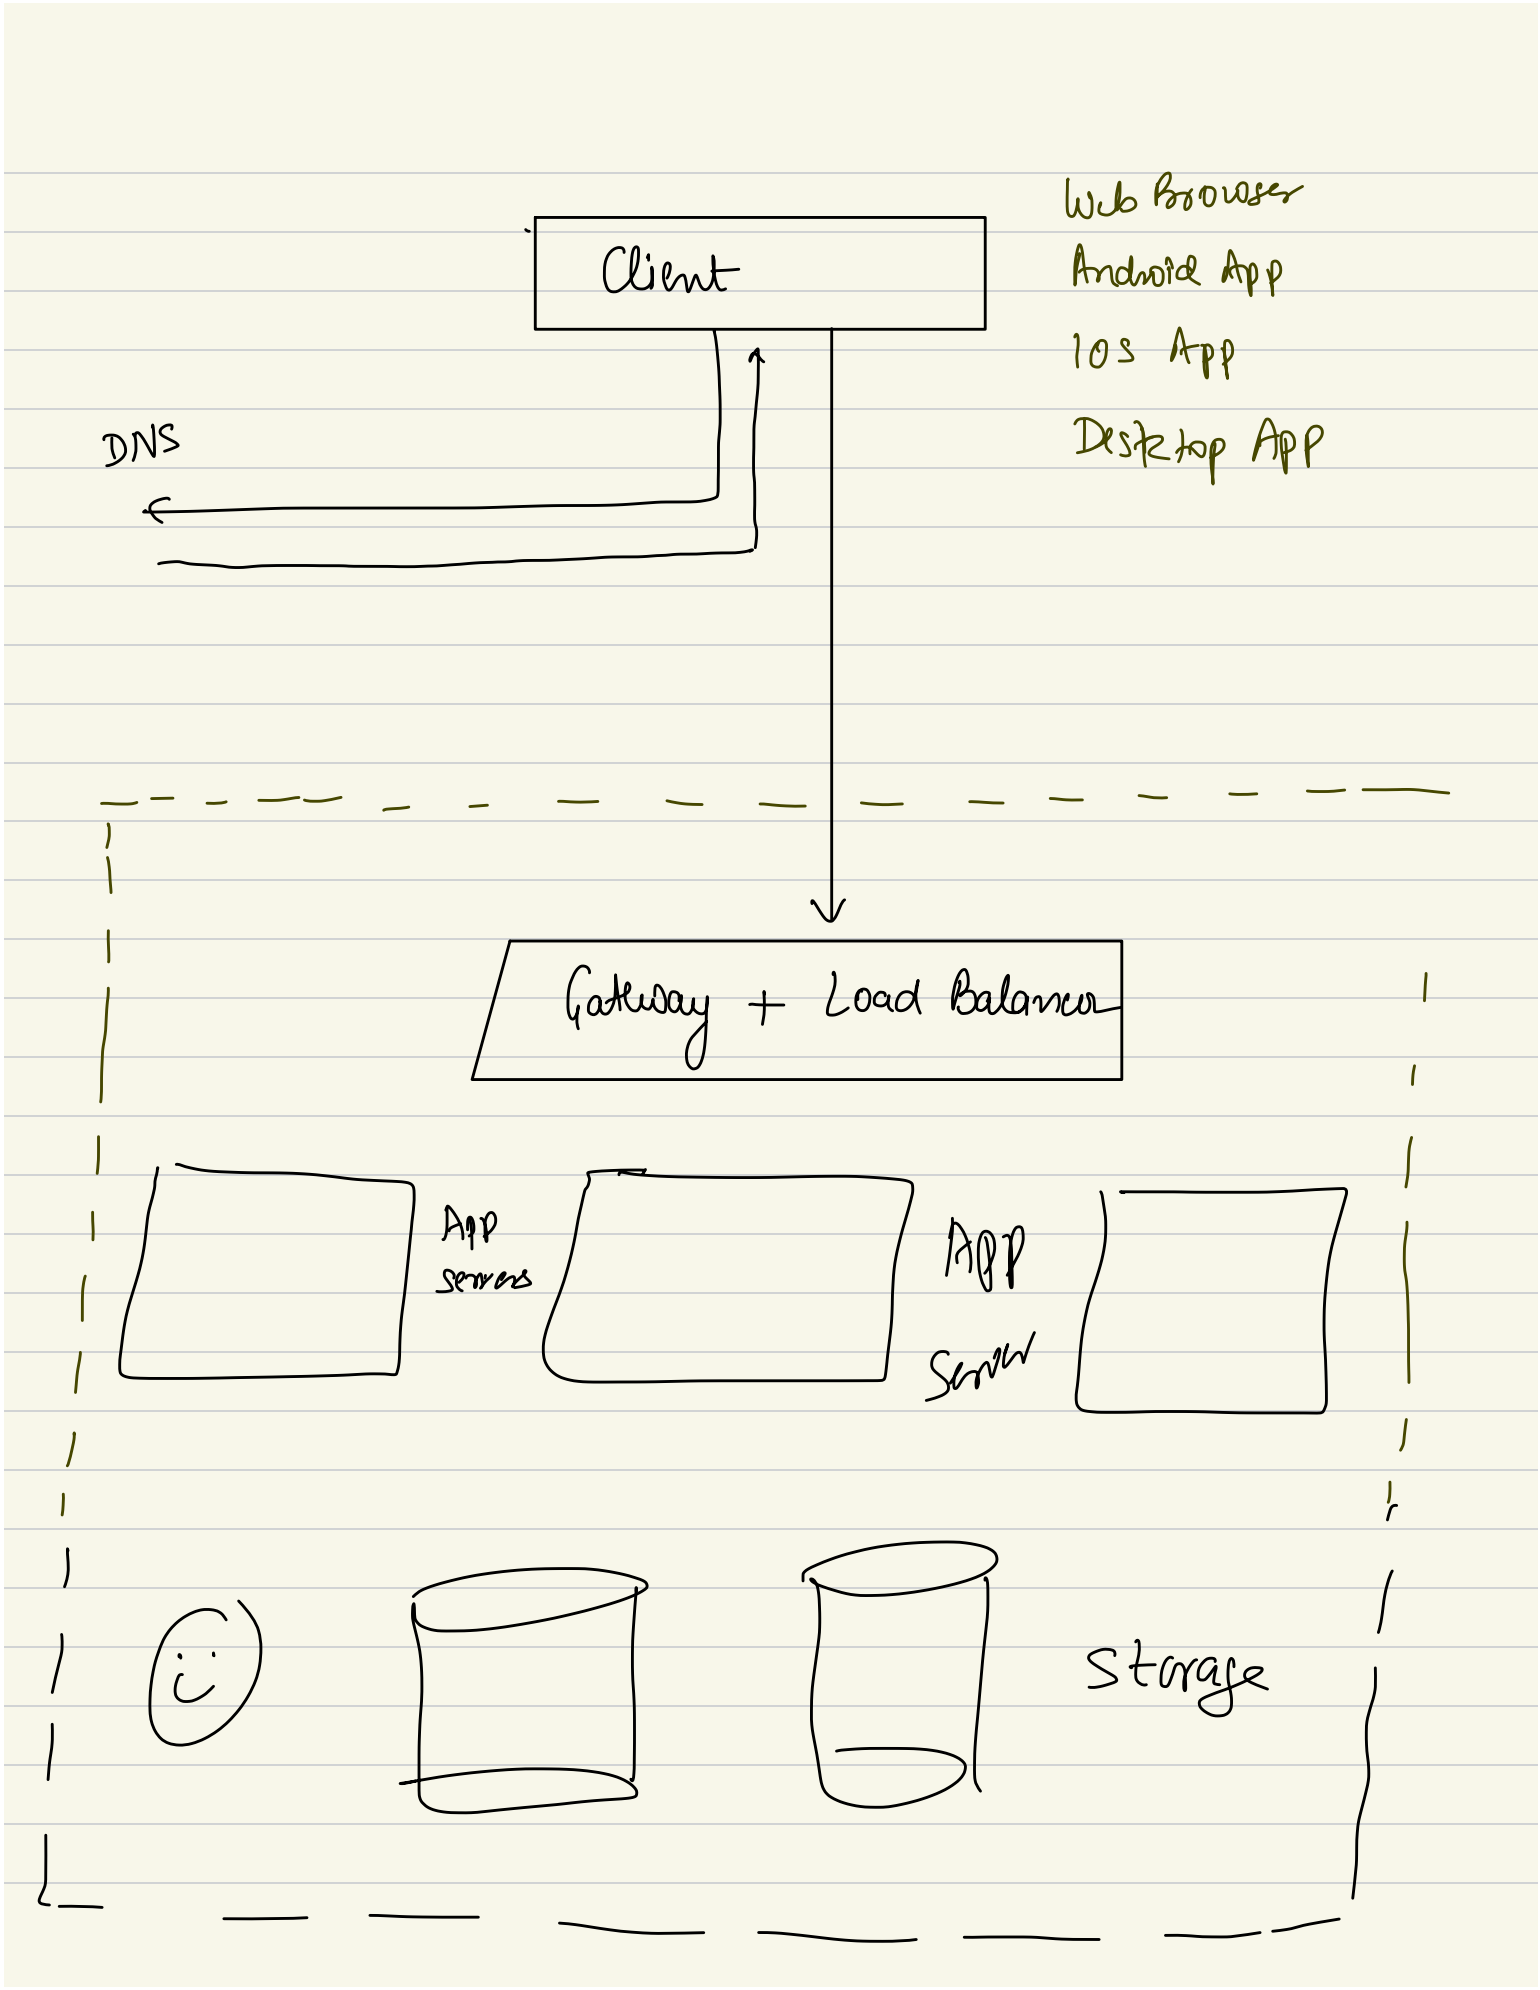
Client
Web Browser
Android App
IOS App
Desktop App
DNS
Gateway + Load Balancer
App
servers
App
Server
Storage

In the backend, we don't just have one kind of machine doing everything for us

[Jack of all, master of none !!! 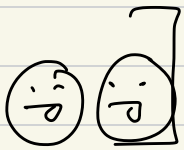]

Instead,

we could want to employ different machines with different resources for different jobs.

For eg.

① compute layer | App Server layer

App Server / compute machines which are going to have good RAM, good CPU so that code and computations can run better.

(2) Storage Layer / DB Layer

HDD ↑↑

SSD ↑↑

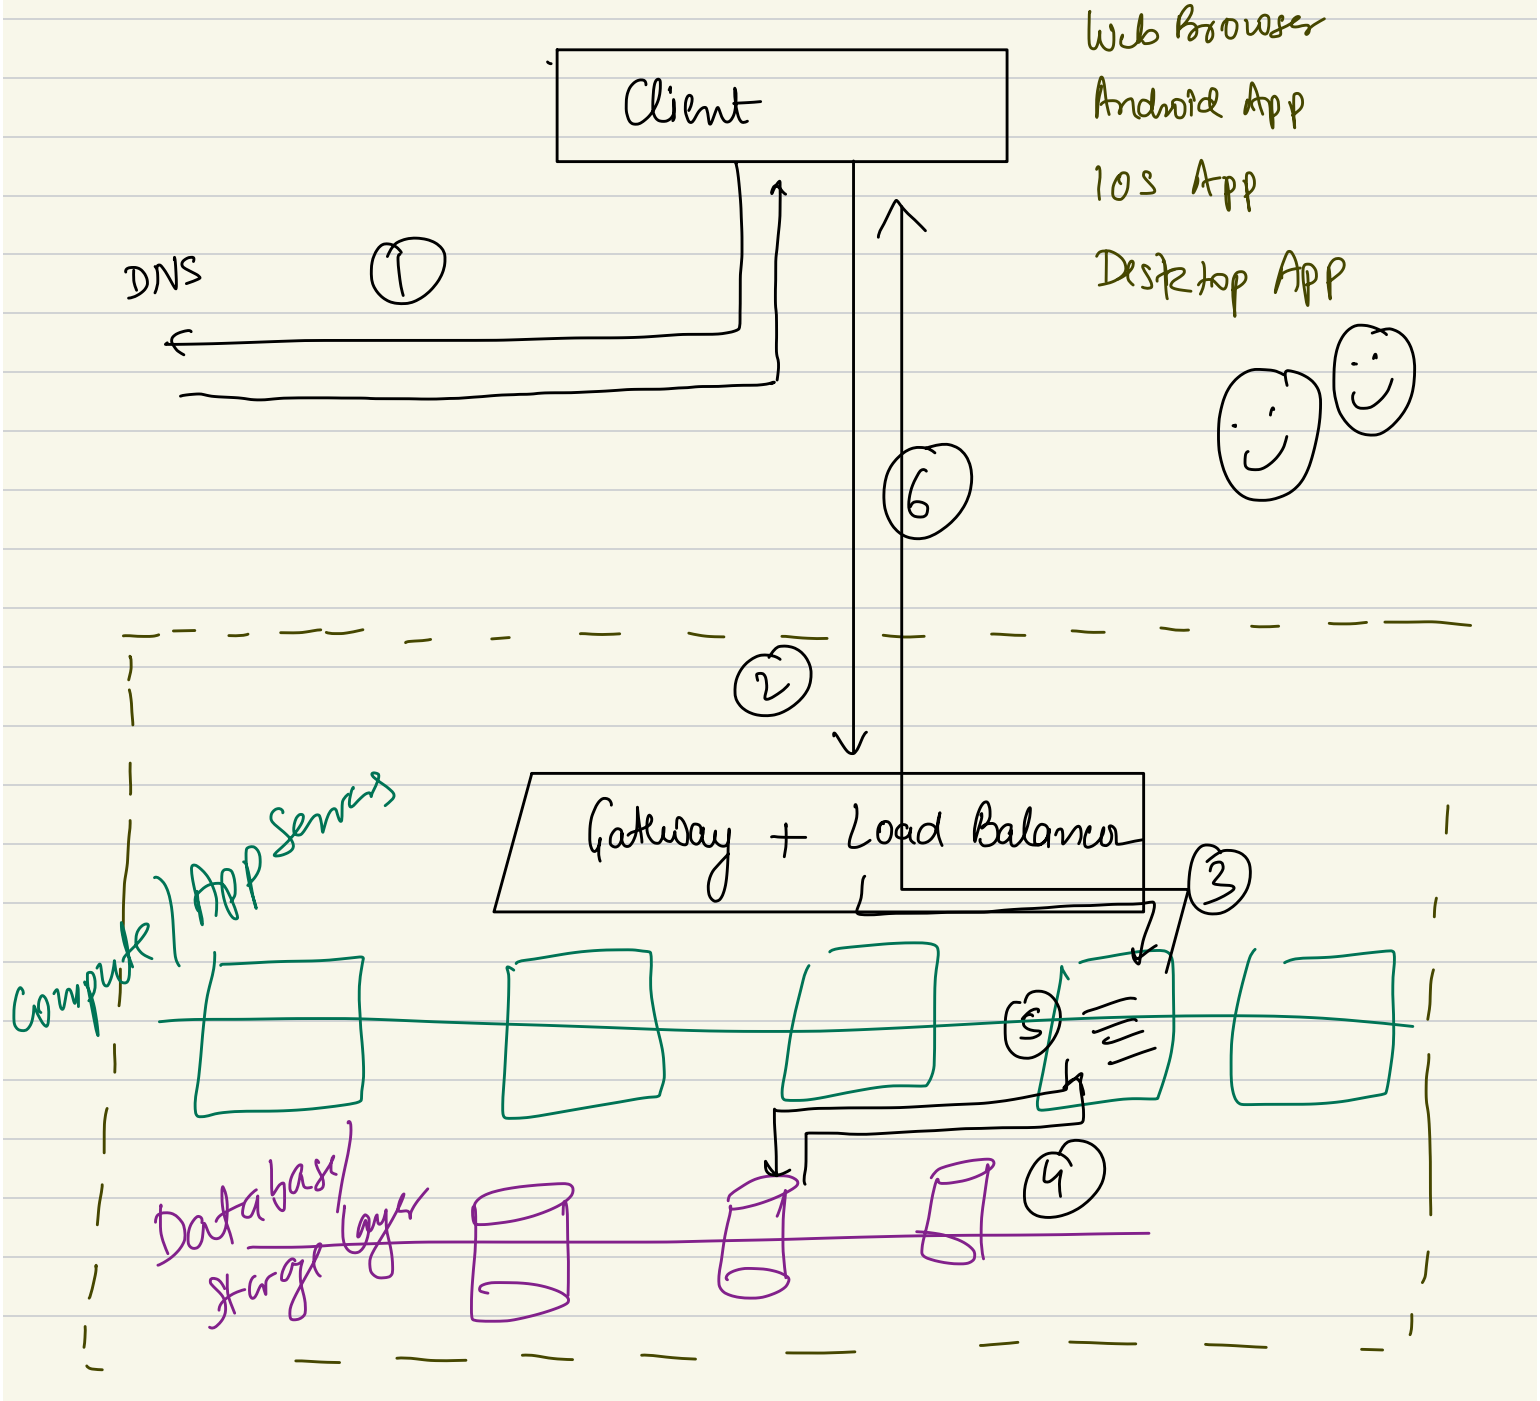

Why decoupling

1. Specialized resources.
2. Single Responsibility Principle
3. Scale your systems better.
4. Code deployments become smoother

Case 1 : Coupled system

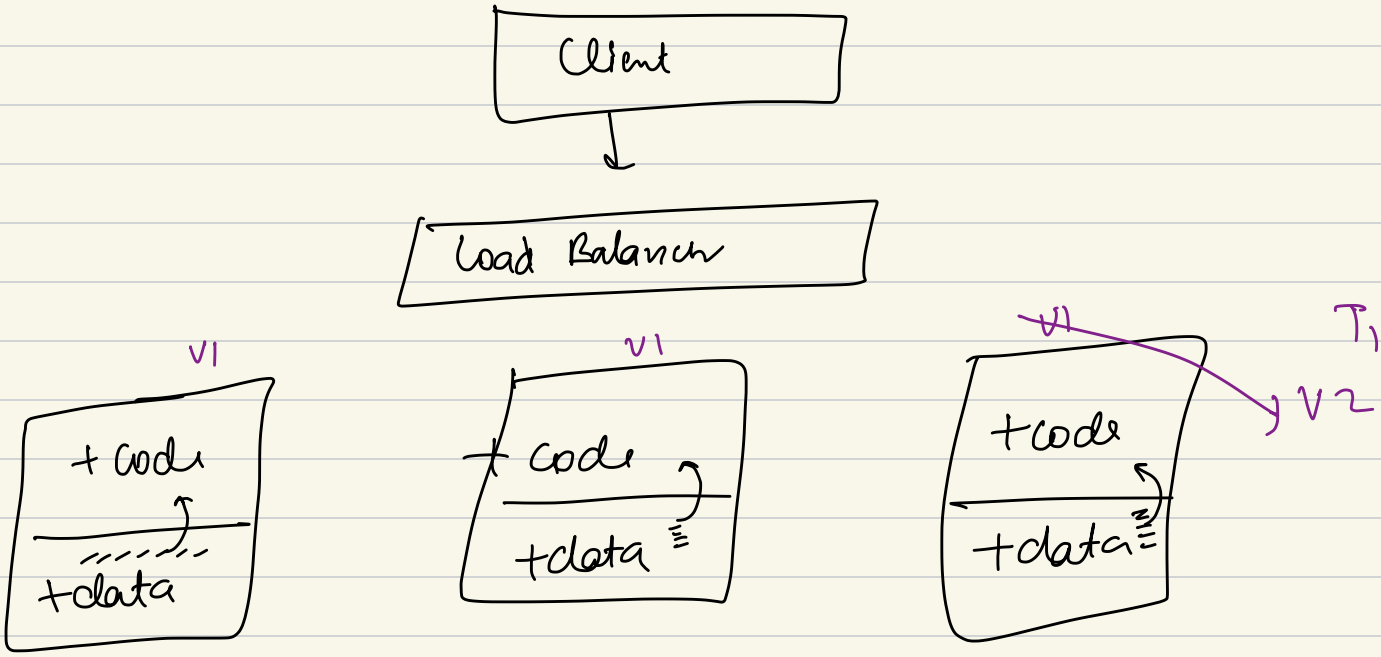

Case 2: decoupling compute machine and storage machine

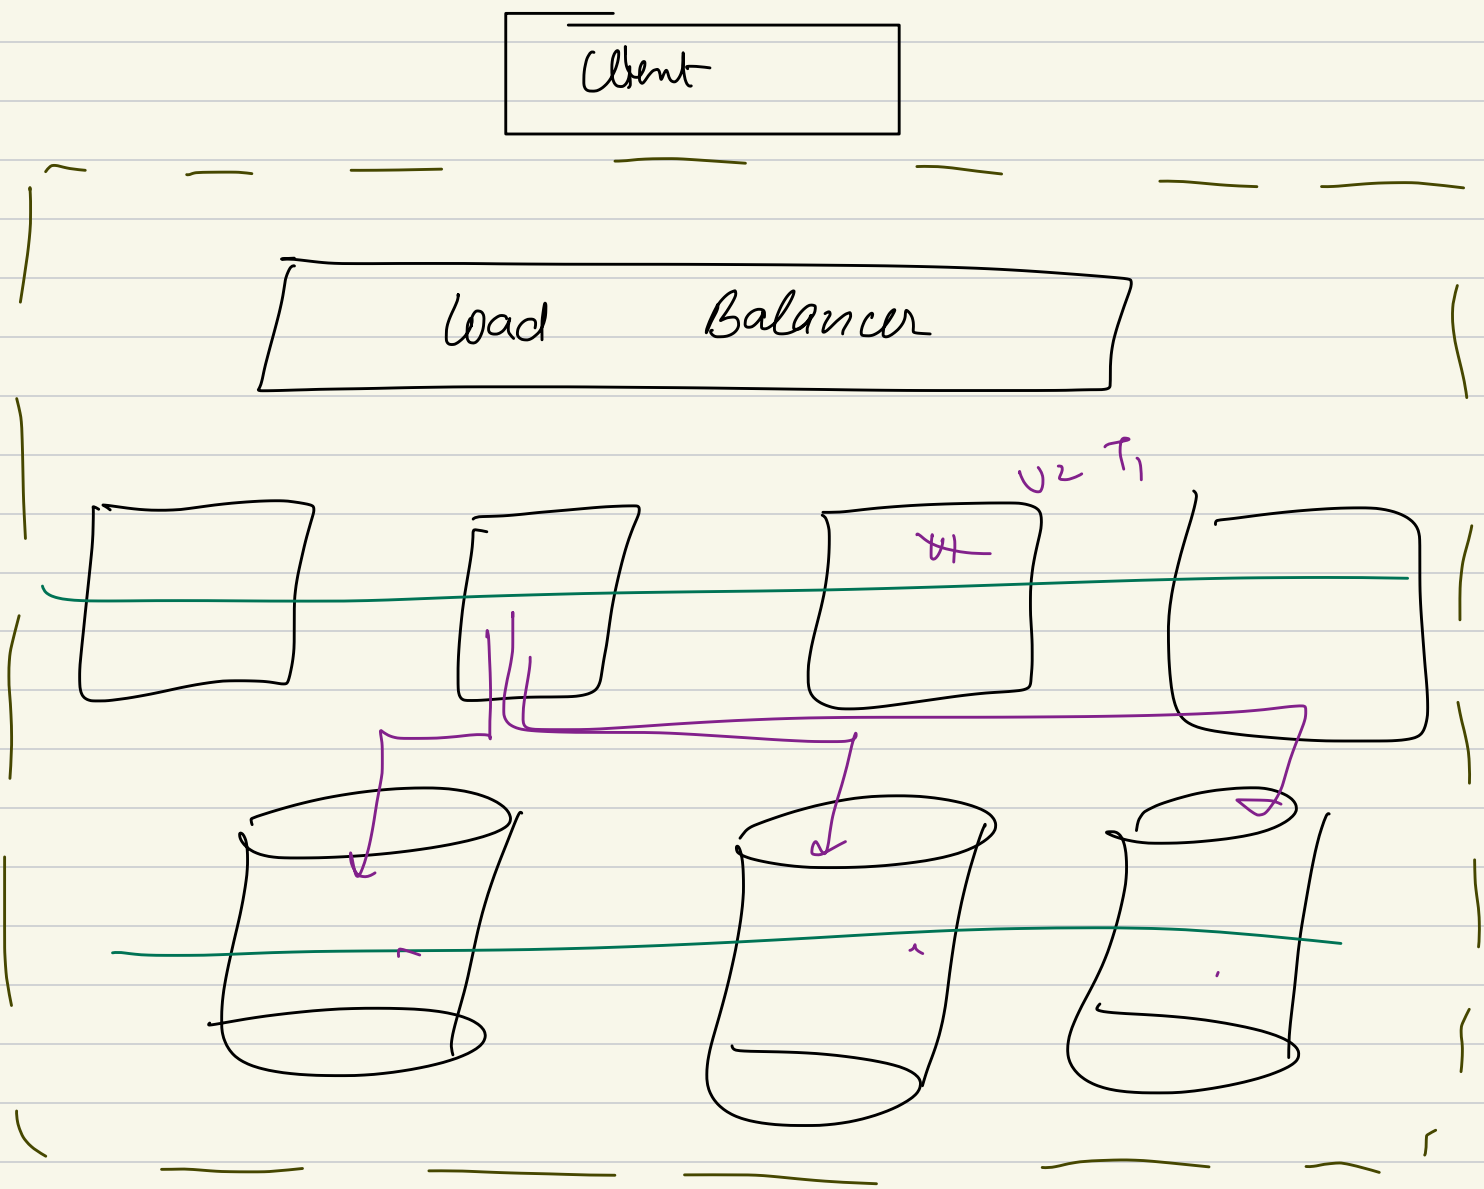

New Reality

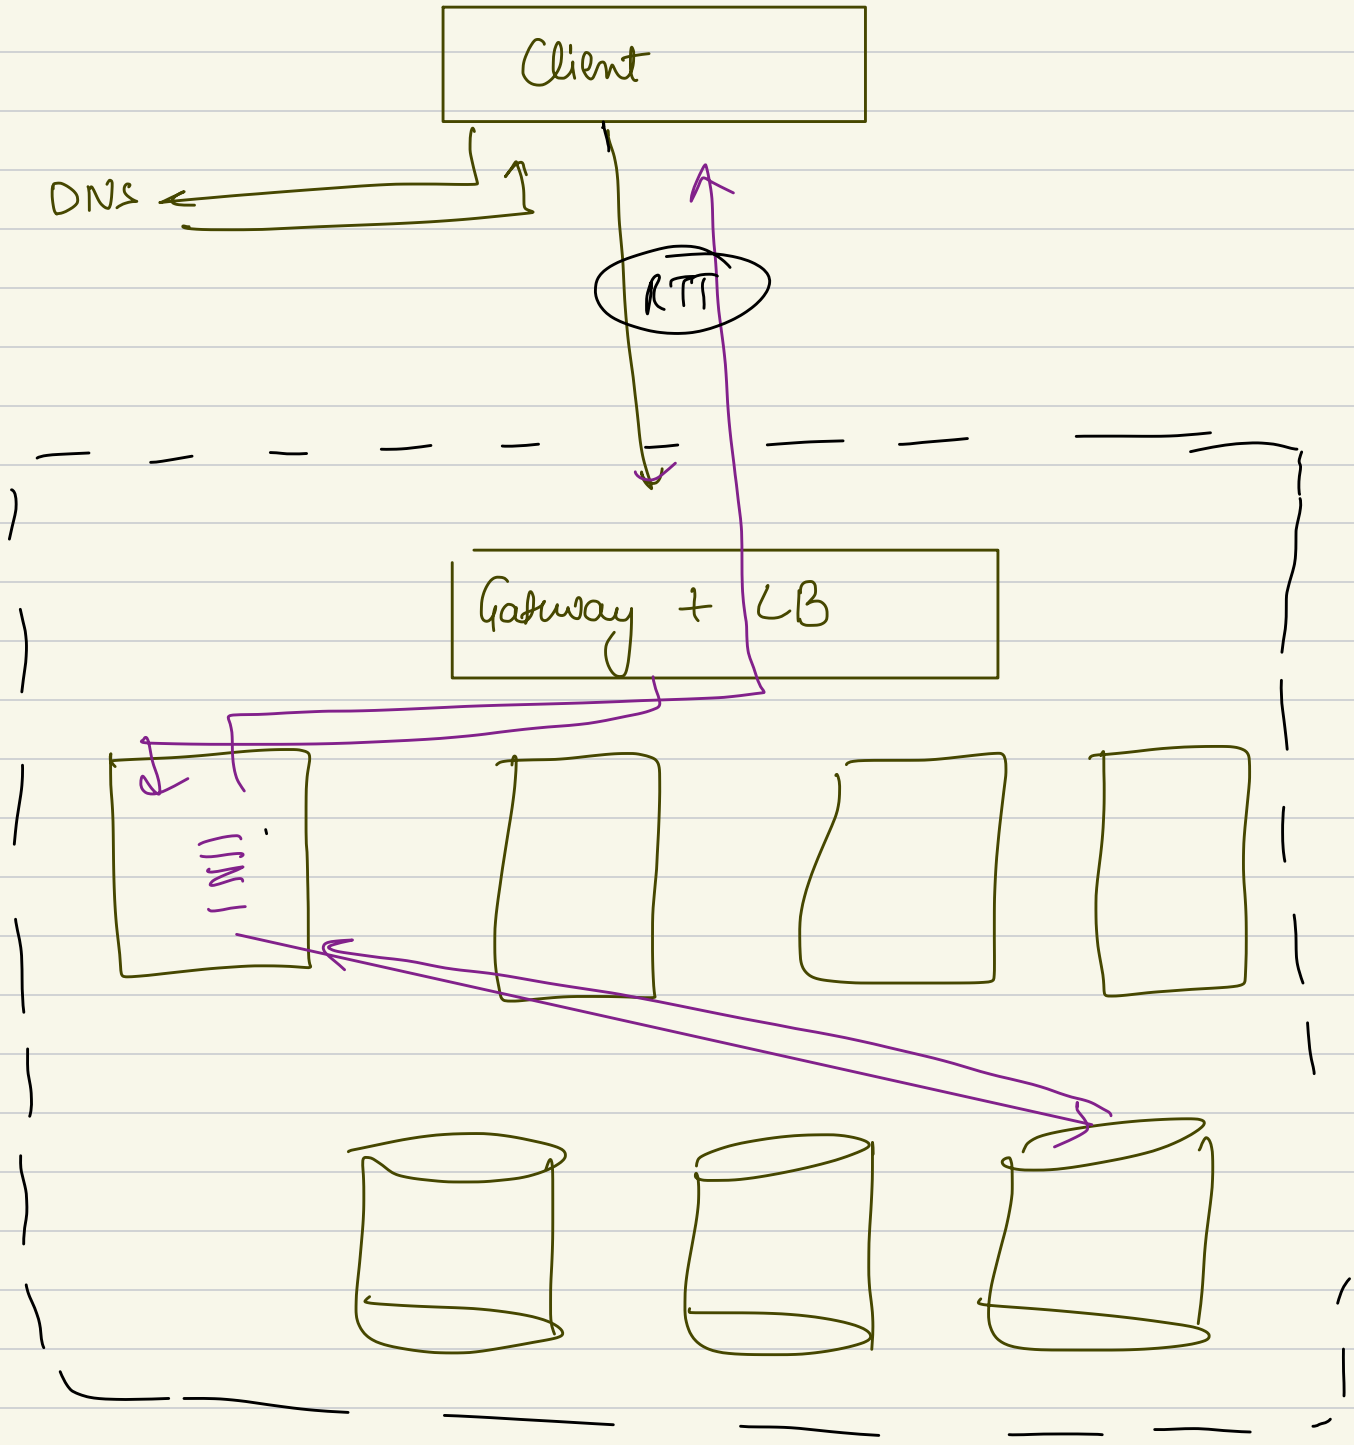

Latency ≡ [Time at which the request got completed − Time at which request started]

[Turnaround Time]

User Latency

= RTT (Round Trip Time)

+

Network Call b/w App Server and DB

+

DB Access Time

+

Computation Time

hops

Caching: The idea of creating partial copies of underlying data so that we can access that data quickly. This is called caching

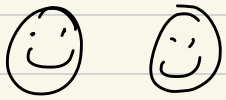

for example in DNS

1. Browser cache
2. OS cache
3. Router cache
4. Neighbourhood DNS cache

(1) Browser | Client Side Cache

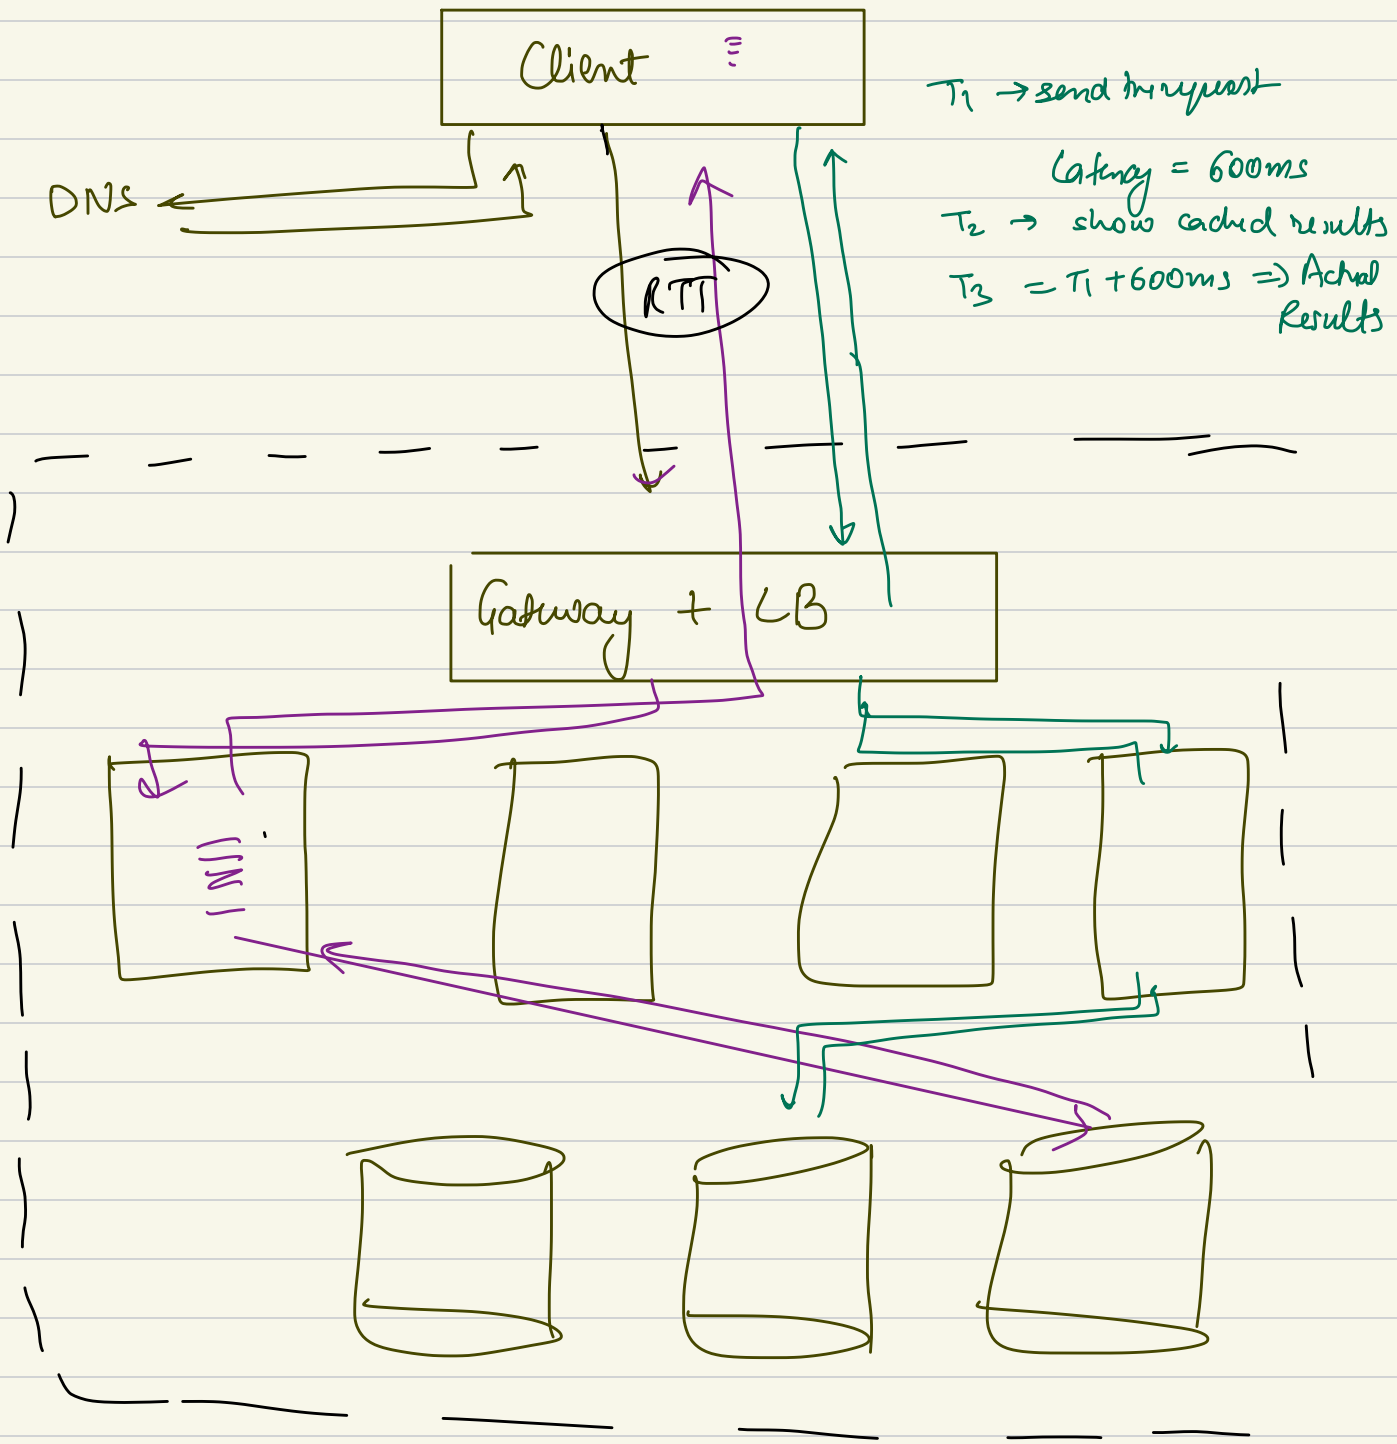

Using client side caching (app data, browser cache)

The client engages the user by showing cached data, so that the perceived latency of the user reduces.

→ "Client Side Optimizations"

② App Server Cache

Fb.com

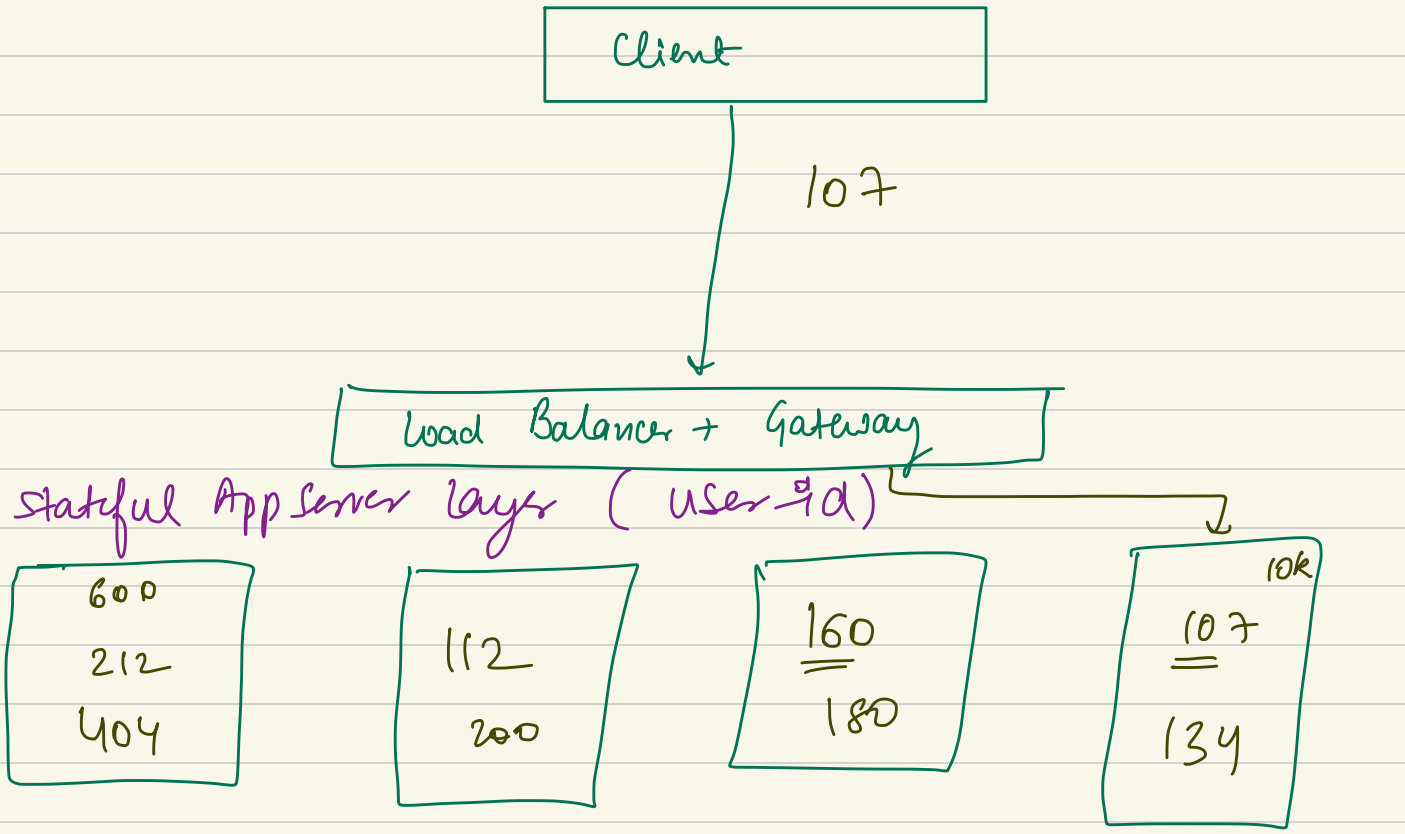

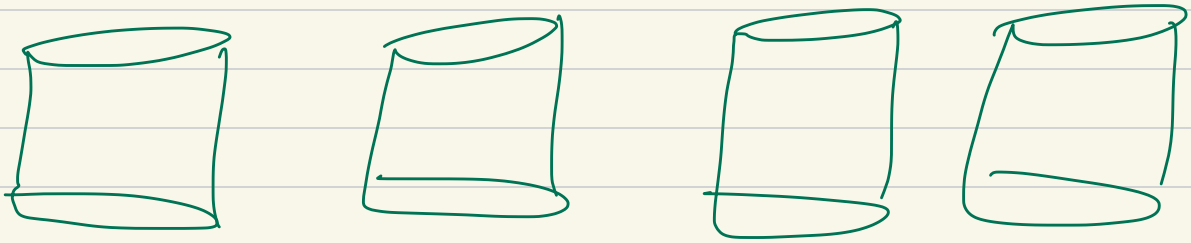

② App Server Caching / Local Caching ☺

Given we know that in stateful load balancing, a particular request is always going to land on the same machine, hence, we can make use of this and cache the relevant data on the relevant machines.

(3) Global Caching

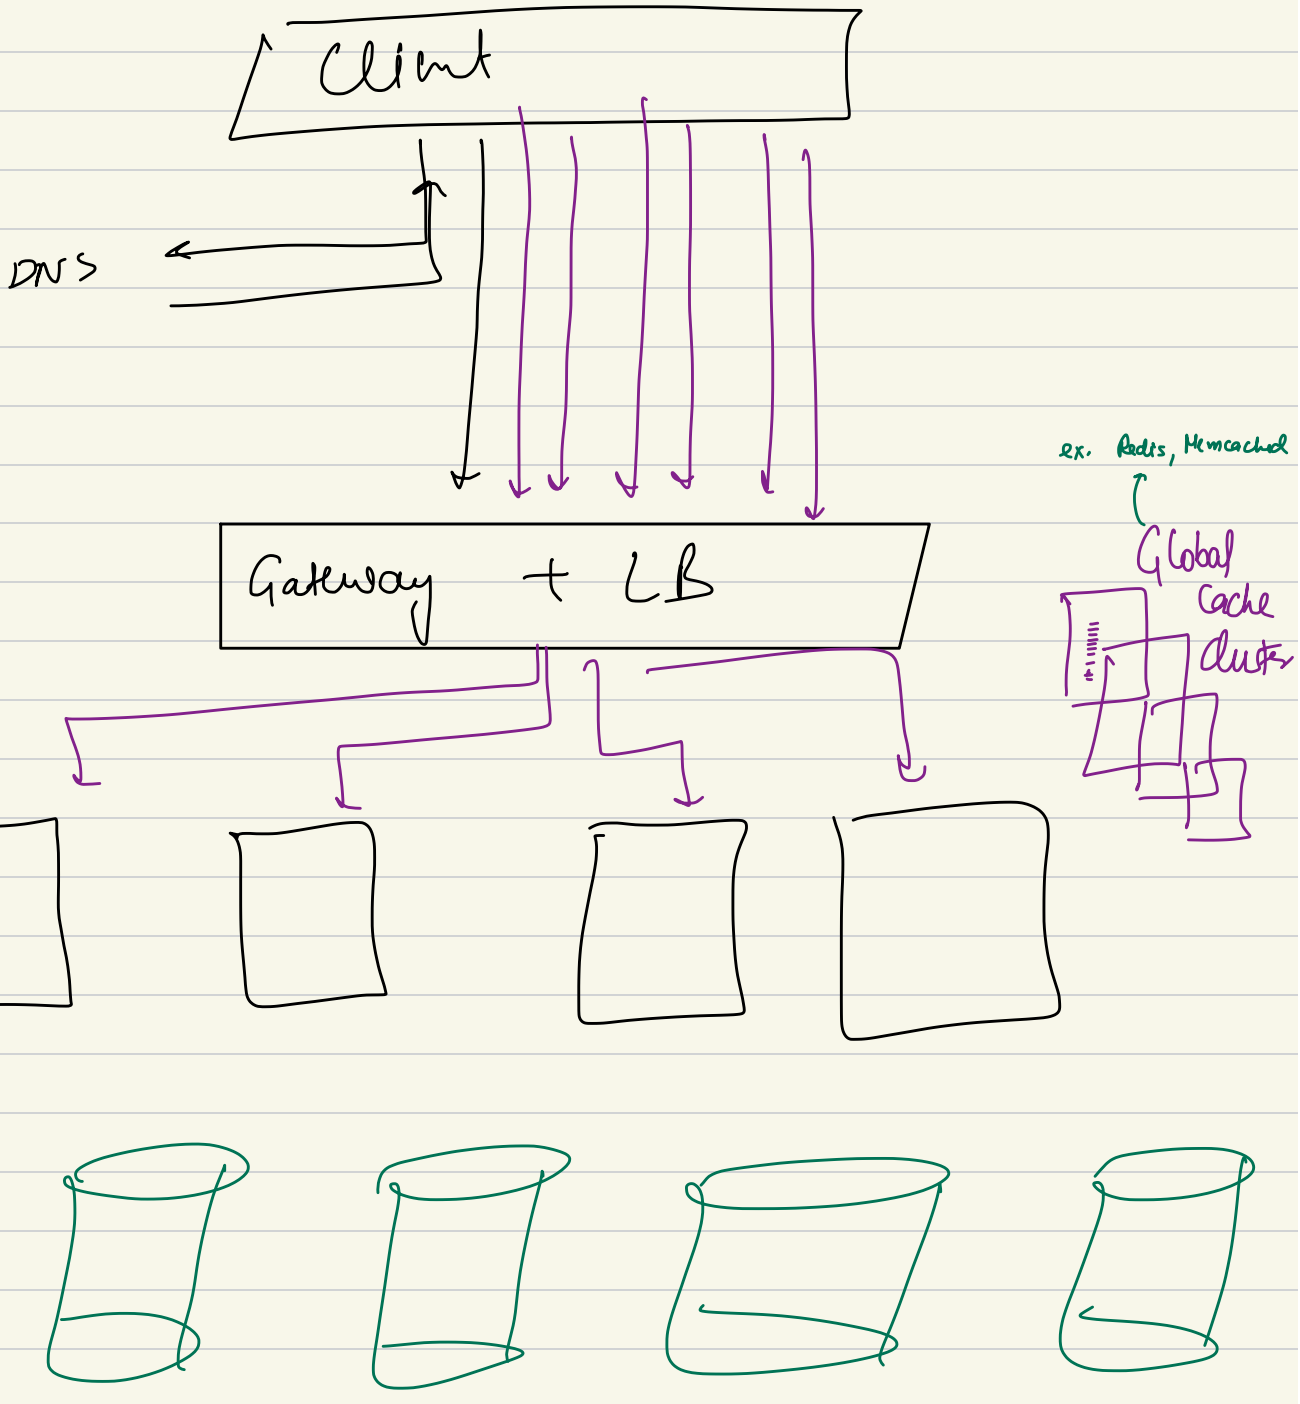

4 CDN

Questions

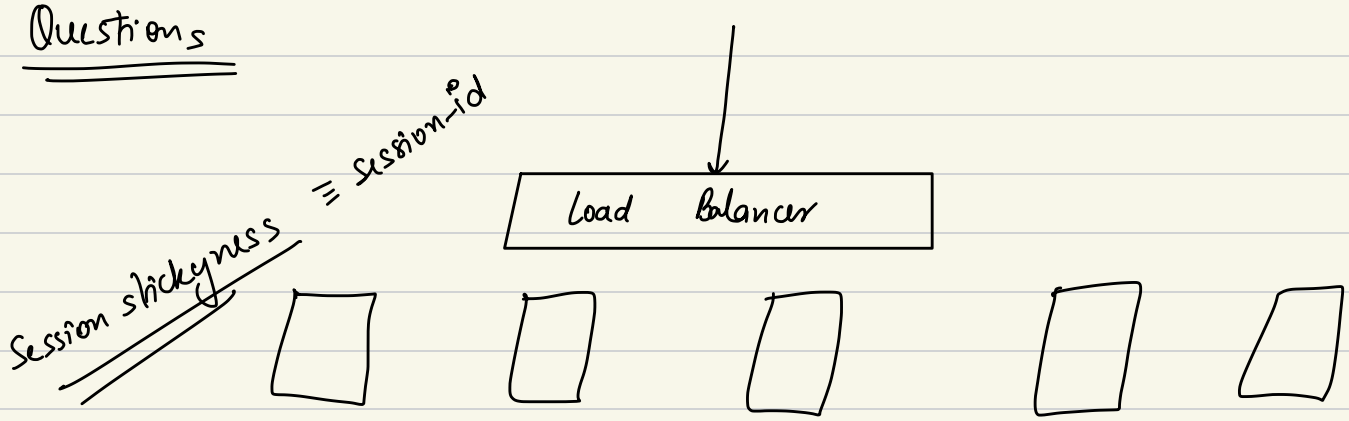

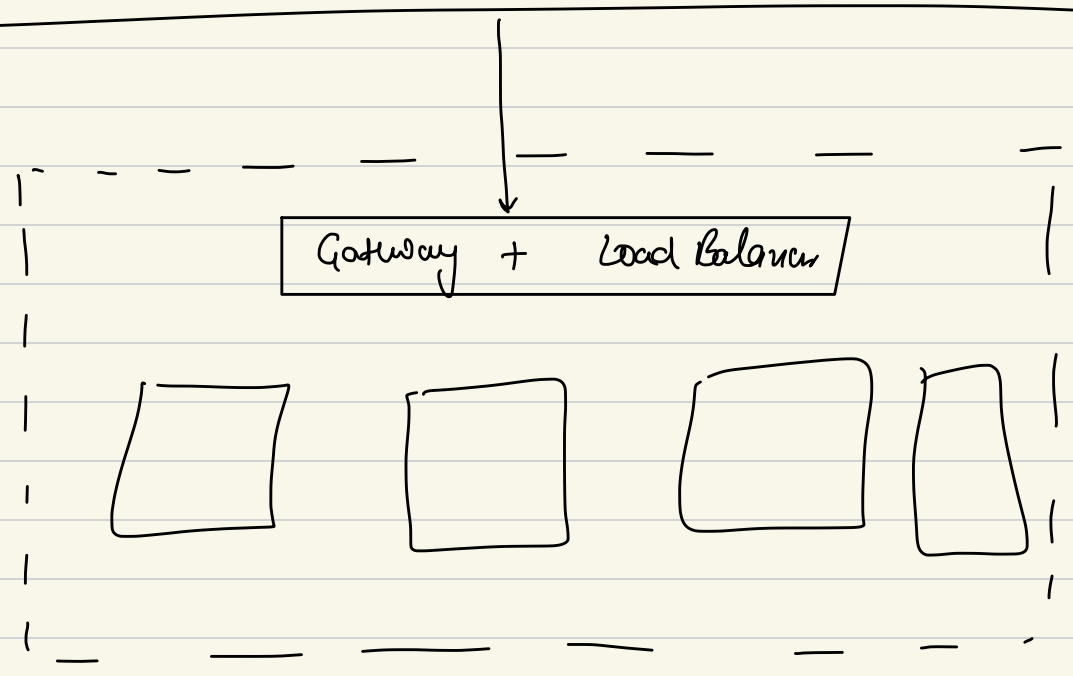

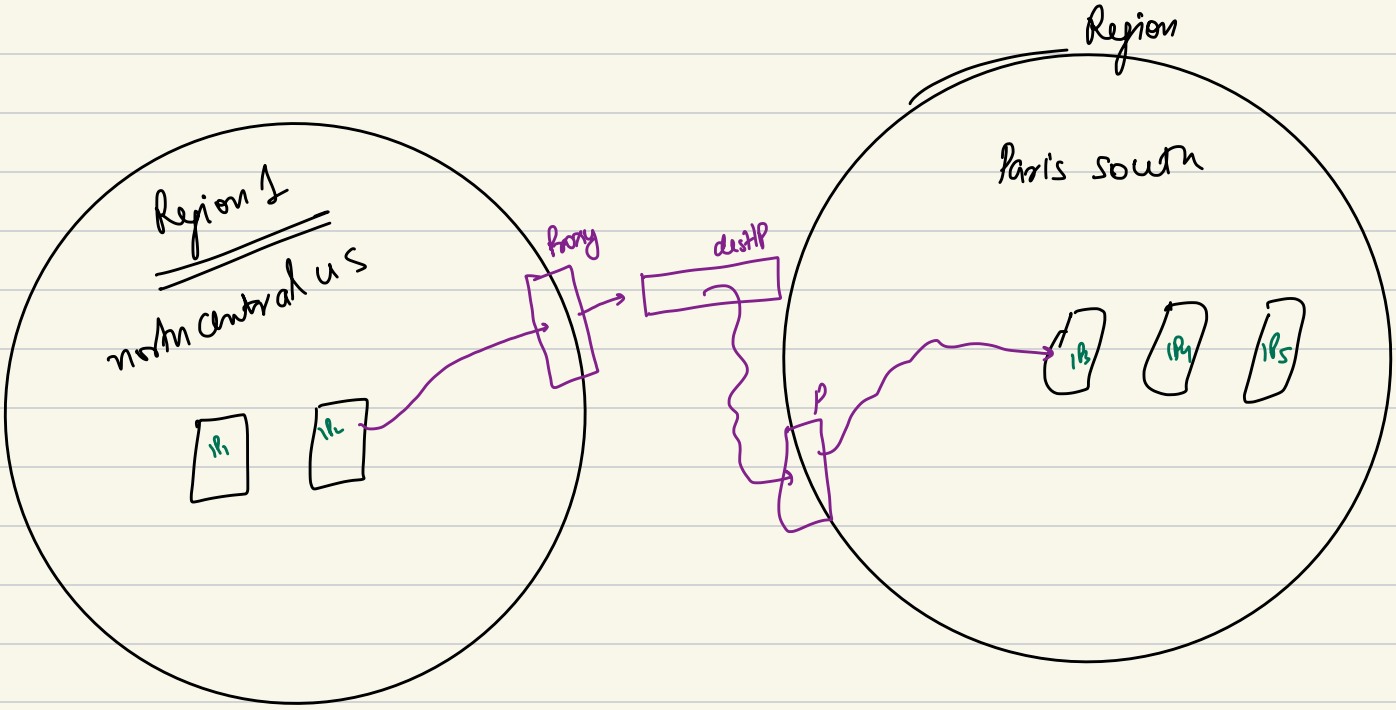

The same consistent hashing logic that we discussed for App Servers, the same is applicable for database layer as well; in fact the same is applicable for the caching cluster as well.