



28/Nov/2023

Facebook Newsfeed + Storage layer concepts







FB


1 Billion

Vertically scaling

Horizontal scaling



→ mutually exclusive
→ collectively exhaustive



1 Billion
/ user data
username
name
last
gender
age
...

1-16



17-50



51-59



60-70



71-85



86-100

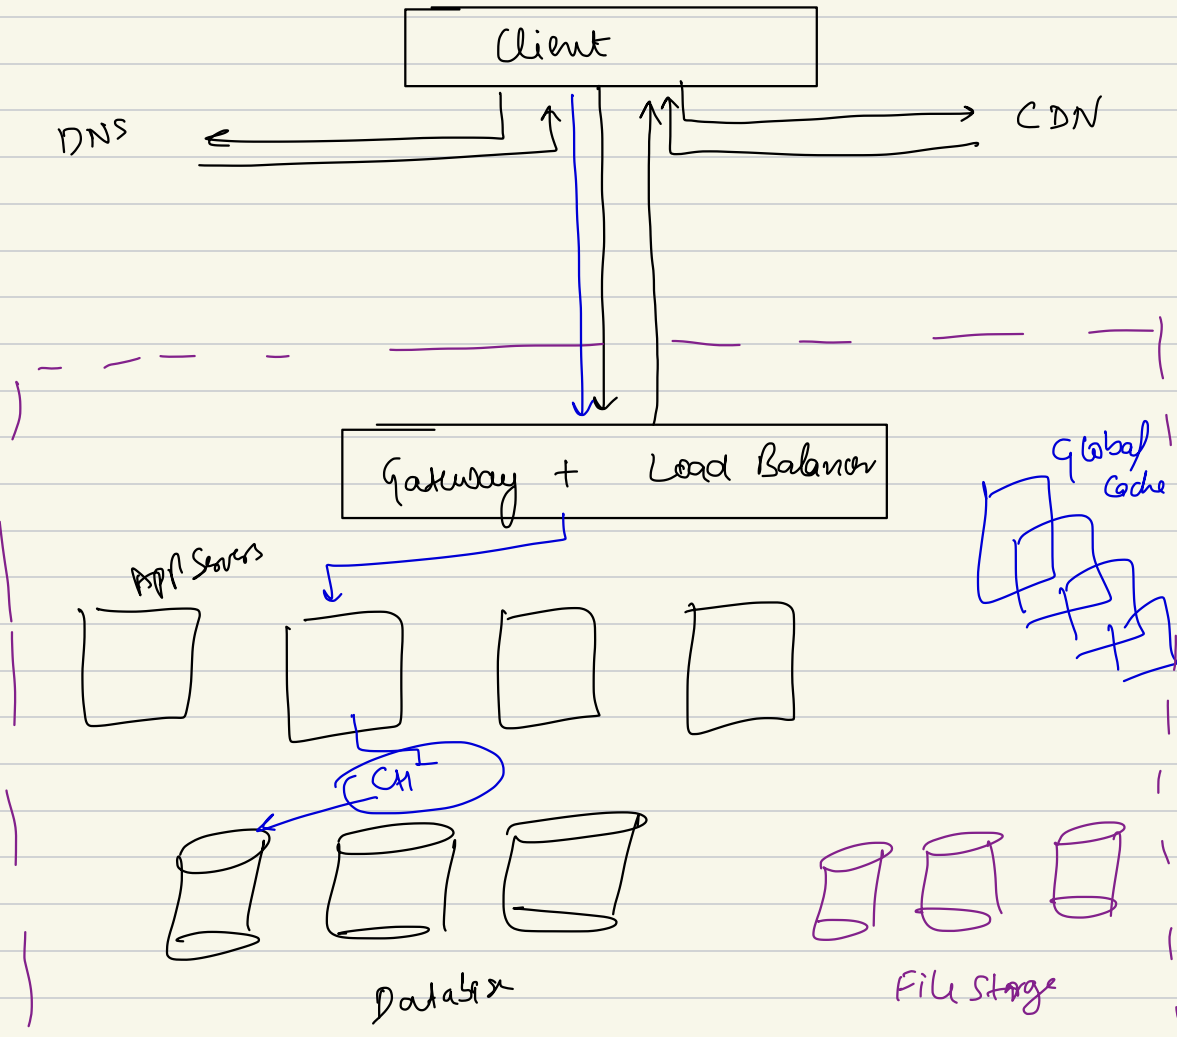


Shards


Shards

Shards

Sharding
||
Partitioning



FB



Mohsin → 100
Deepak, Virhal ↓ 250
200

user				
id	name	email	gender	last seen
pr				


user-friends

user-id	friend-user-id
100	200
100	250 :
200	100
250	100

Posts

id	text	userid	timestamp	...


CM



sharding-key \rightarrow request-id
OR

hashing-key


\downarrow
partition key



very optimal
😊


INTRA SHARD
QUERY

detour



Bad 😞

INTER SHARD
QUERIES



Intra Shard queries



Inter-Shard Queries



We need to choose the right sharding key so that majority of our frequent queries can reduce as intra shard queries, and only few rare queries for not time critical

queries can be handled as intershard.

<u>1 shard</u> (<u>intra</u>)	<u>2 shards</u> -----	<u>All shards</u> (<u>inter</u>)
Best	Good enough	Worst 😞

10:16pm - 10:25pm 😊

Facebook Newsfeed : Case Study

✓ (A) Newsfeed → Posts by friends / pages you follow

✓ (B) Profile Page → Posts that you yourself made

Assume : Facebook uses SQL databases.

User				
id	name	email	gender	last-seen
par				

user-friends	
user-id	friend-user-id
100	200
100	250
200	100
250	100

Posts				
id	text	userid	timestamp	...

Assumption: All data of FB can fit inside
the same machine.

Profile Page

> Posts made by you ^{user-id} X

```
select *  
FROM  
POSTS P  
WHERE  
P.user-id = X
```

LIMIT 10
OFFSET 50

Newsfeed


> Posts made by the friends
of that user with user-id = X

```
SELECT *  
FROM  
USER-FRIENDS F JOIN  
POSTS P  
ON F.user-id = X  
and F.friend-id = P.user-id  
AND  
P.timestamp < Now() - 30days  
LIMIT 10 OFFSET 0
```

Pagination 😊


Reality: All data can't fit inside
single machine 😊 😊

Consistent hashing
Shard-Key



user-id \equiv sharding key

this above decision means the following


- (1)
 - (a) user-attributes
 - (b) list of friends
 - (c) posts made by that user
- (2) 

Profile Page

> Posts made by you

user-id
x

```
select *  
FROM  
POSTS P  
WHERE  
P.user-id = x
```

intra shard


LIMIT 10
OFFSET 50

Newsfeed


> Posts made by the friends
of that user with user-id = x

```
SELECT *  
FROM  
USER-FRIENDS F JOIN  
POSTS P  
ON F.user-id = x  
And F.friend-id = P.user-id  
AND  
P.timestamp < Now() - 30days  
LIMIT 10 OFFSET 0
```

All shards going


P

0




debut

Denormalization

Redundancy

optimal queries ☺☺☺


100 — 107



Calculations

Users : 2 Billion

DAU : 500 Million



(80-20) Rule

users → post / day = 5 Million

10 posts/day ⇒ 50 Million posts / day

post	post-id	→	8 bytes
	user-id	→	8 B
	timestamp	→	8 B
	text	→	40 B
	video / image location	→	20 B
	tag	→	20 B

= 104 B → 400 B


$$\underline{50\text{ M}} \times \underline{400\text{ B}}$$

$$20 \times 10^3 \text{ MB}$$

$$\boxed{\text{Post Data / day} = 20\text{GB / day}}$$

$$20\text{GB} \times 30 \text{ day}$$


$$= 600\text{GB of posts} \text{ 😊 😊}$$




Commodity hardware

Storage commodity hardware \rightarrow ^{-8TB} 4TB of data / machine


Data Retention — 10 years





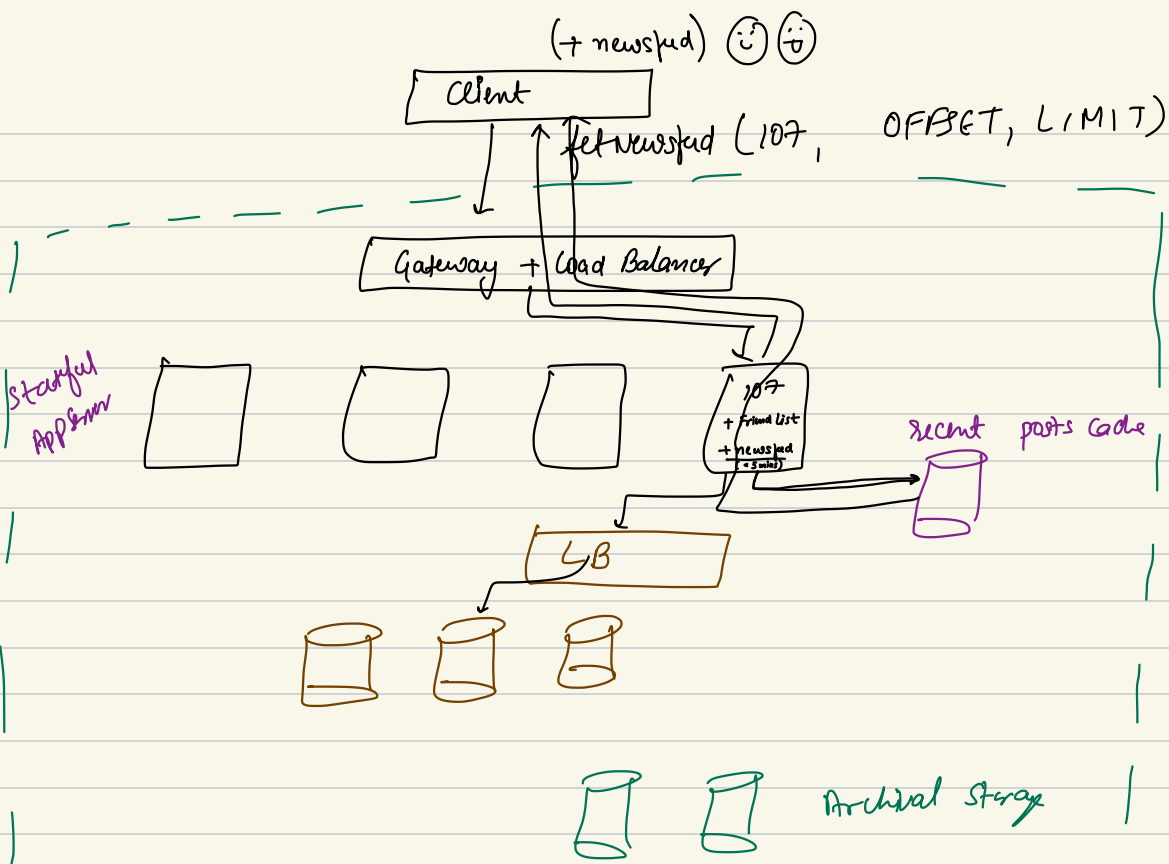
Load Balancer



Using a recent posts cache we can converted the inter shard query of newest generation and converted it into a intra shard query. 😊







Question

For a user 240 user-id \Rightarrow combine all posts that they liked + assign 200

102



111

