

20/12/2023

Inverted Index -

Apache Lucene

+scalability

Elastic Search

doc:
A


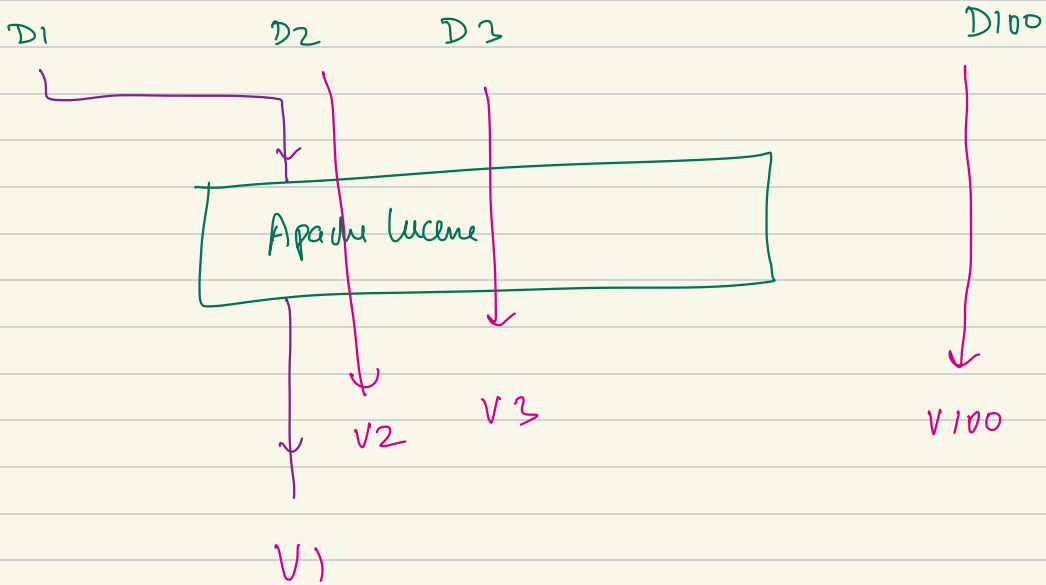
India is a beautiful country.
It is an emerging world power.

~~India is a beautiful country.~~
~~It is an emerging world power.~~

beauty
emerge

	emerge	Score
India	→	x
<u>Beauty</u>	→	y
<u>country</u>	→	z
emerge		
world	→	x
power		
India beauty		
beauty country		

100 documents that I want to index



Inverted Index

Superset of key \leftarrow union of all vectors

A Post:

India is a beautiful country

B Post:

Canada India relationship is at
a new low

C Post:

Canada engineering colleges might
or shom.

D Post:

India lost the CWC

Post:

I love my job. I love my life

ex. Unigrams.

Inverted Index

India - A (0.4), B (0.2); D (0.4)

beauty - A (0.3)

country - A (0.3)

Canada - B (0.4); C (0.3)

relation - B ()

new - B

low - B

engineer - C

college - C

sham - C

loose -

CWC -


love -

job -


my -

life -

query
→
"Indian relations"




query;
vector(query)




A, B, D

B




example: 10k documents at Google



query

vector(query)

Search on Inverted Index



[list of documents]

[sorted list of documents] ★

Apache Lucene is great. It does great
NLP, fast query processing 😊

But Lucene is NOT scalable.

Elastic Search

— Lucene made scalable 😊

③ Design Tradeoffs

① Availability >>> Consistency
😊 😊




② ✓ low latency for getting Search Results
★ Low Read latency
✓ we are fine with medium-high latency for writes...

④

Design deep dive

① . compute resources .

② . storage .






① sharding is definitely needed

② Replication as well.

data you're storing

- ① documents $D_i \rightarrow 10B$
- ② vector $V_i \rightarrow 10B$
- ③ Inverted Index $\rightarrow 1$

CN \rightarrow doc-id




This means that we now have a
distributed inverted index 😊 😊

WRITE Flow (a document ingestion)

ingest a new document

new post \equiv new document



doc204 \Rightarrow

"India is the best"

vec - 204 \Rightarrow


India - 0.6

best - 0.2

India best - 0.4

Logical Shards and Physical Nodes / Machines


Traditional




Logical Node = Physical Machine

2 concepts a play.

① Physical Machine




② logical idea of shard → Master slave etc




10 shards, each with 2 slaves


1- 10 shards




30 logical machines




physical cluster of 5 machines






15 logical shards



Elastic Search exploit the above idea

~~100 machines~~ in ES cluster $\rightarrow +10$ physical machines



ES

✓ 100 physical machines

✓ 70 Master machines | shards

3 slaves

280 logical shards

logical \rightarrow 280



$\frac{280 \text{ logical}}{100 \text{ physical}}$



$\frac{280 \text{ logical}}{110 \text{ physical}}$



We have used different logical shard and physical machine concept.

Because every logical shard has some documents allocated and its respective inverted index.

We don't want to create/destroy our logical shards

as creating / destroying them would mean

re allocation of documents in turn meaning

re creation of inverted indices

This is time consuming....

So we want to avoid this.

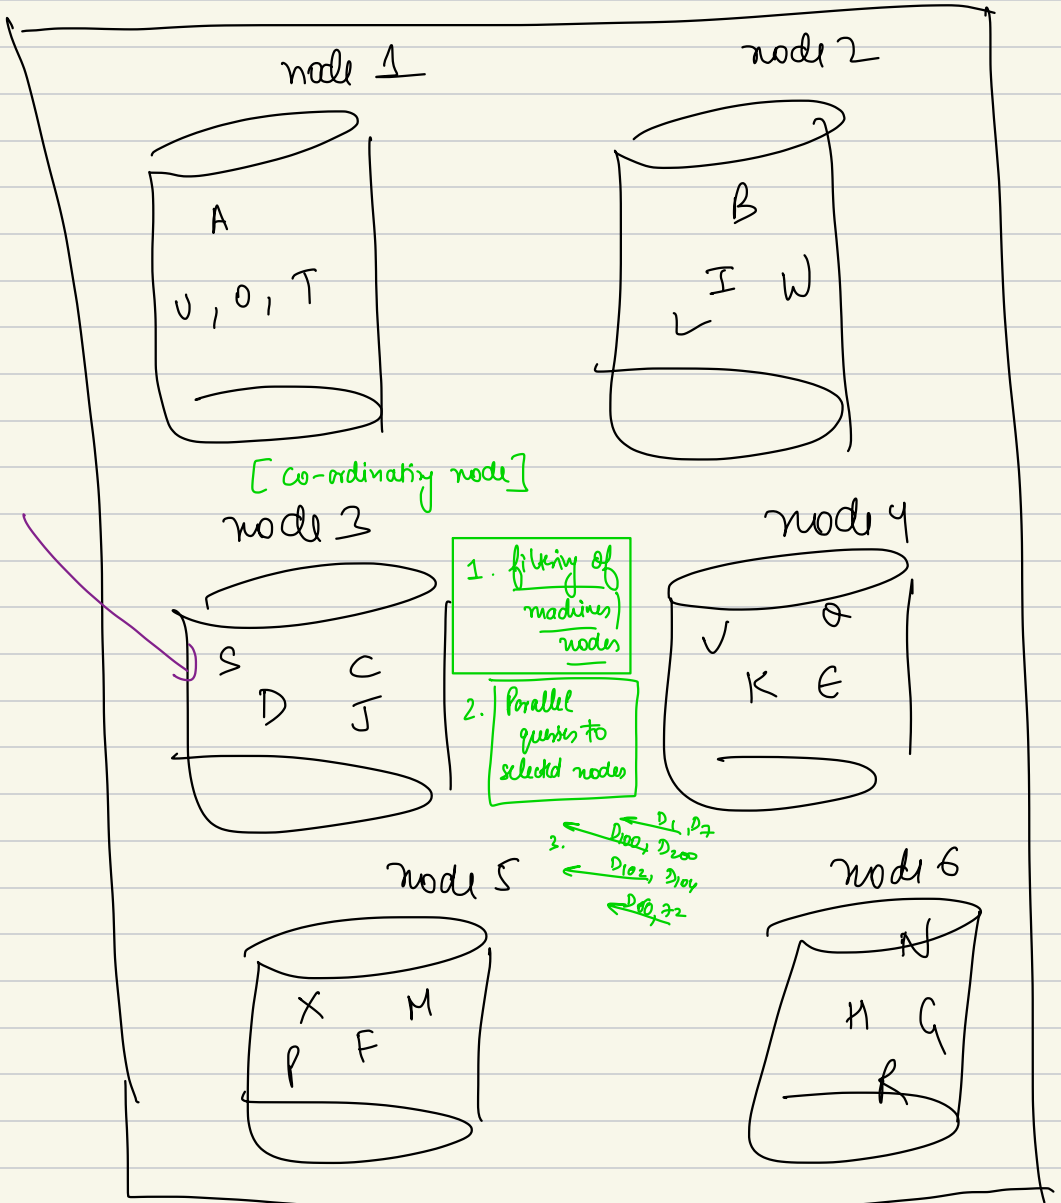
Rather, we can add new physical machines, keeping the logical shards as it is, so that

we can pick up entire logical shards


and put them to new machines

without having the need to recreate indices.

Read Flow



8 masters



8 masters

256ms/master



24 logical shards

query $\xrightarrow{\text{computation}}$

raw query + vector(query)

Search engine for Forbes.com

- ✓ forbes.com/spats/cricket/indvspak-cwc-match.
- ✓ forbes.com/spats/cricket/ind vs aus-cwc 2003-final
- ✓✓ forbes.com/spats/tennis/wimbledon-records
- ✓ forbes.com/cars/tesla-latest-model



Co-ORDINATING Node

① query processing \rightarrow $\begin{matrix} \text{vector of query} \\ \left[\begin{array}{c} - : s \\ - : q \\ - : r \\ - : q \\ - : s \\ - : w \end{array} \right] \end{matrix}$

\longrightarrow land at a co-ordinating node






② [Optional]

Filter out some master shards as important shards based on what your query contains.


[Step ② is only possible if the sharding of documents happened in a logical clustered way / topic based partitioning]

③ Out of the non-rejected shards, the coordinating node will decide a list of physical machines which have


Representation from each of selected M-S logical shards.




I



II



III



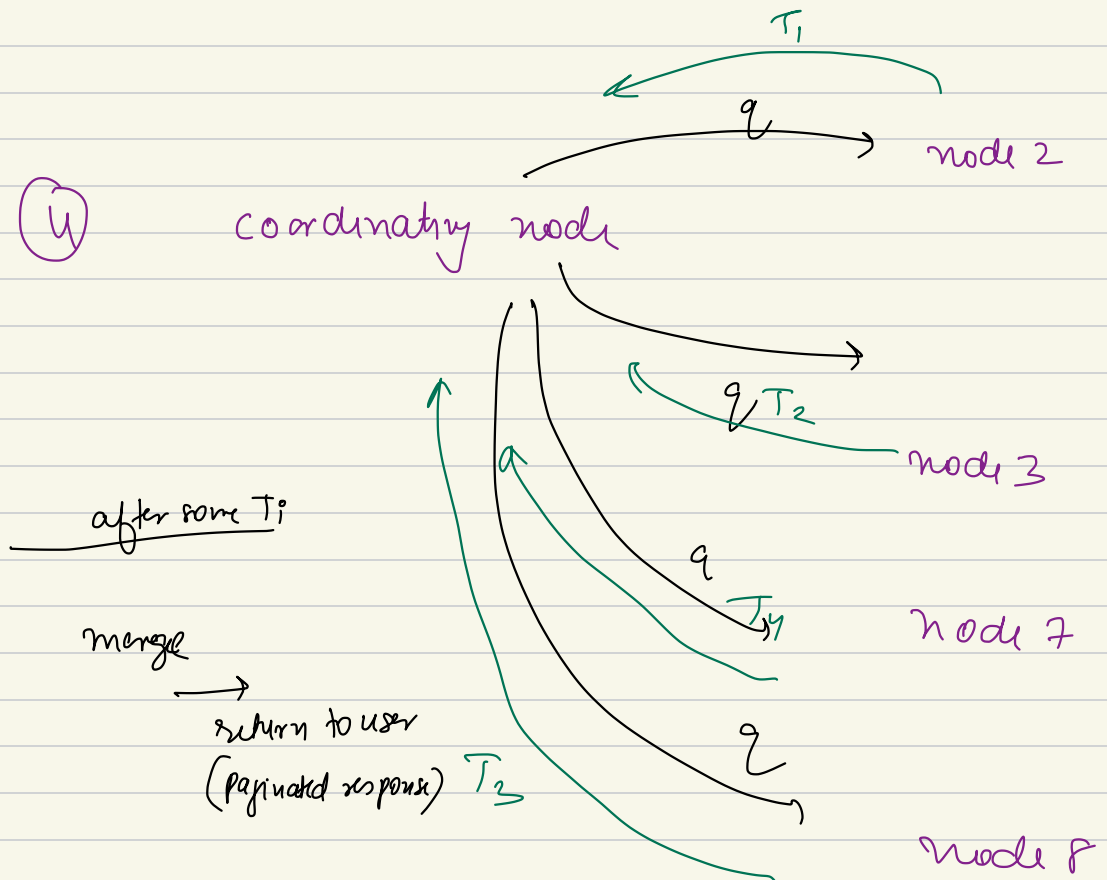
IV

You will select such a list so that
number of physical machines you
pick up on is minimum

+

load of machines is low

In the shard selection, the coordinaty node will identify some physical nodes (least number of machines) and parallelly send queries to selected machines.



Parallel Processing of Query

Timeout

(combine + rank)



User



①

S3



File/Blob Store



ND FS




②


Quad Tree



User



border shops near me



Questions

①

	Score
<u>India</u> -	0.4
Country -	
india test -	
india pm -	
Modi -	0.5
Narendra Modi -	
Australia -	0.1

NLP Techniques

① Frequency of word



②

TF-IDF



③



Q



+ 1 Machine



A B C